

CS 6110 Software Correctness, Spring 2022 (edited from a previous class; look for bugs!)

Ganesh Gopalakrishnan
School of Computing
University of Utah
Salt Lake City, UT 84112

URL: bit.ly/cs6110s22



Slides for Lec4 : Agenda

- Logic readings
 - Discussions
- FSM material
 - Not familiar to many
 - Not taught as a practical tool in most courses
 - Not taught as the first real formal methods in most courses
 - E.g. context-free parsing is really the first dramatically successful FV tool
 - The horrible bugs in them
 - Fortran missing comma -> spacecraft in a coma ☺
 - The)))*(((rule for parsing!
 - Knuth's article on how bad compilers were in the 1960s !!
 - Youtube version is pretty good (how he got his honeymoon paid for!)
 - https://archive.computerhistory.org/resources/text/Oral_History/Knuth_Don_1/Knuth_Don.oral_history.2007.102658053_all.pdf
 - OpSem basics → Sutter

Slides for Lec4 : Agenda

- How the project discussions are proceeding
 - Notes being taken
 - References being catalogued
 - Not all are at the same point in learning about Logic, Automata, ...
 - But they are all important to learn properly!
- We will go through the Bradley/Manna book Ch 1,2
 - Discussion of readings from Ch1
- SPIN Usage
 - Long trace
 - Search depth control

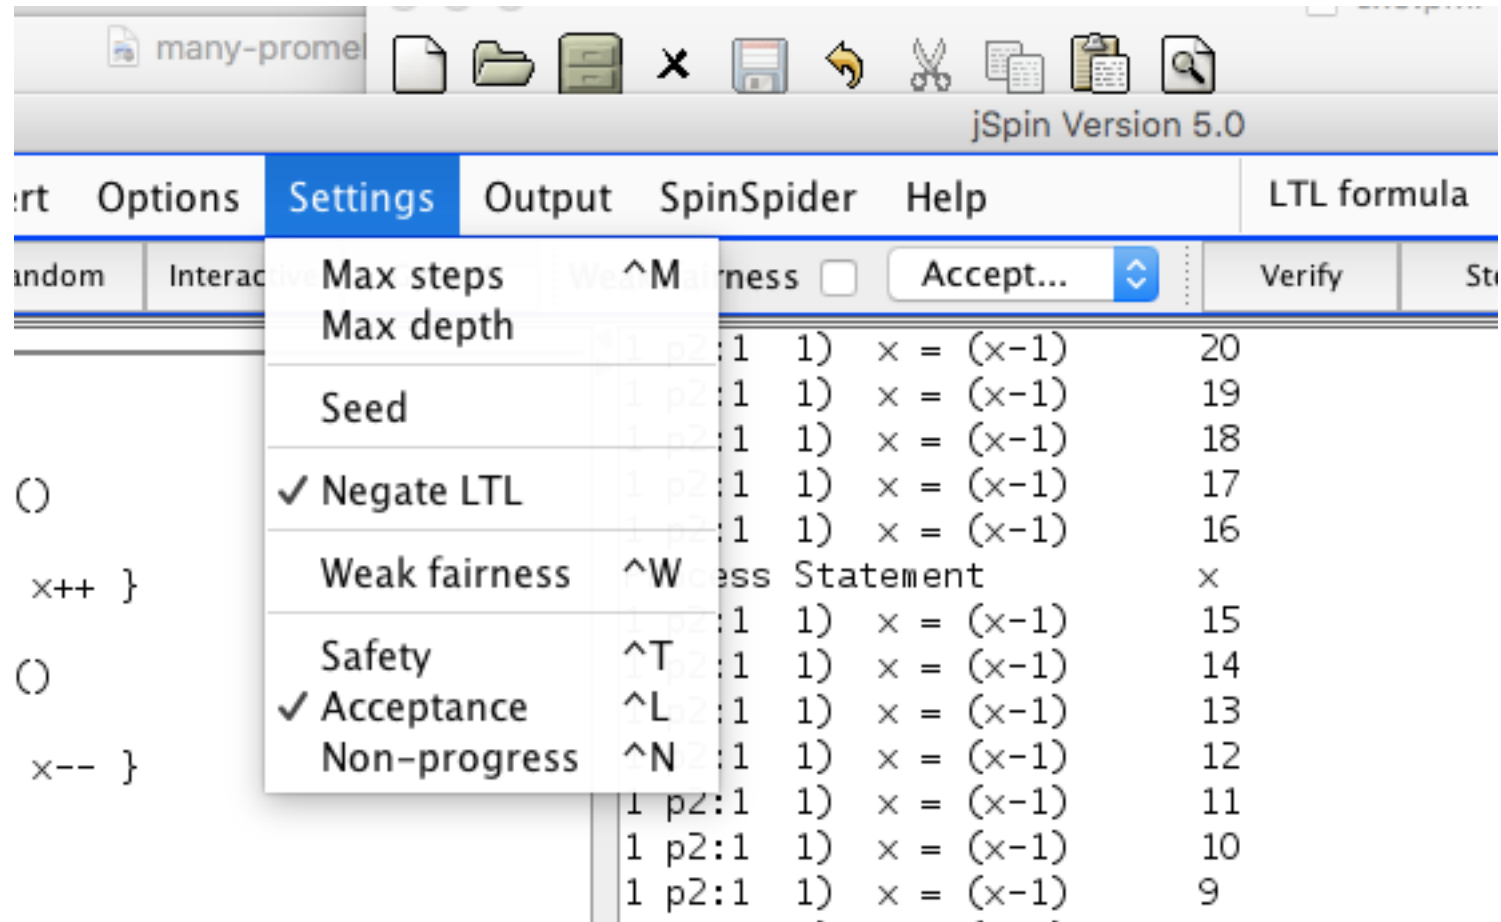
Slides for Lec4 : Agenda

- Readings for next week + Asg-2
 - Posted tomorrow (1/21) evening
 - From Ben-Ari's book
 - From CEATL
 - About NBA \neq DBA
 - About LTL
 - Nested DFS
 - How we can show LTL validity by creating the right Kripke Structure!
 - Reinforces idea of validity
- Distributed Locking Protocol
 - In Promela
 - In Murphi

How SPIN helps

- Helps you quickly model “situations”
 - See SPIN-Soldiers + SpinTutorial by Ruys
- Debug your thoughts
- Find flaws in design
- Has a lot of power and will NEVER become irrelevant
 - The more complex a system gets, the less you need to get fancy wrt verification
- (after a 1-month training), you can do this
 - When in doubt, whip up a SPIN model and find bugs (in system or in your head)
- Has worked for large systems
 - Recent correctness workshop -> entire new Pthreads implementation modeled!
- We don't have the full month to train you
 - This is the last SPIN week ... but I'll give you a REAL protocol verif to read on Thu!

ex5



ex5

The screenshot shows the jSpin Version 5.0 application window. The main editor displays a PML file named `ex5.pml / ex4.pml` with the following code:

```
1 byte x;  
2 active proctype p1()  
3 {  
4   do  
5     :: atomic { x++ ; x++ }  
6   od }  
7 active proctype p2()  
8 {  
9   do  
10    :: atomic { x-- ; x-- }  
11  od }  
12 never {  
13   do  
14     :: x==4 -> break  
15   od  
16 }
```

The **Settings** menu is open, showing options: **Max steps**, **Max depth**, **Seed**, ☒ **Negate LTL**, **Weak fairness**, **Safety**, ☒ **Acceptance**, and **Non-progress**. The **Acceptance** option is selected.

The right pane shows the execution trace:

```
1 p2:1 1) x = (x-1) 20  
1 p2:1 1) x = (x-1) 19  
1 p2:1 1) x = (x-1) 18  
1 p2:1 1) x = (x-1) 17  
1 p2:1 1) x = (x-1) 16  
1 p2:1 1) x = (x-1) 15  
1 p2:1 1) x = (x-1) 14  
1 p2:1 1) x = (x-1) 13  
1 p2:1 1) x = (x-1) 12  
1 p2:1 1) x = (x-1) 11  
1 p2:1 1) x = (x-1) 10  
1 p2:1 1) x = (x-1) 9  
1 p2:1 1) x = (x-1) 8  
1 p2:1 1) x = (x-1) 7  
1 p2:1 1) x = (x-1) 6  
1 p2:1 1) x = (x-1) 5  
Never claim moves to line 14 [((x==4))]  
0 p1:1 1) x = (x+1) 4  
0 p1:1 1) x = (x+1) 5  
spin: trail ends after 254 steps  
#processes: 2  
254:      proc 1 (p2:1) ex5.pml:8 (state 4)  
254:      proc 0 (p1:1) ex5.pml:4 (state 4)  
MSC: ~G line 16  
254:      proc - (never_0:1) ex5.pml:16 (state 7)  
2 processes created  
Exit-Status 0
```

The bottom terminal window shows the following commands and output:

```
/usr/bin/gcc -o pan pan.c ... done!  
/Users/ganesh/software/jspin/jspin/jspin-examples/manyexs/pan -a -m20000 -X ... done!  
/usr/local/bin/spin -g -l -p -r -s -t -X -u250 ex5.pml ... done!  
/usr/local/bin/spin -g -l -p -r -s -t -X -u250 ex5.pml ... done!  
/usr/local/bin/spin -g -l -p -r -s -t -X -u2500 ex5.pml ... done!
```

ex6

```
/* Smart reductions saved states */
byte x,y;
active proctype p1()
{do
  :: atomic { x++ ; x++ }
od
}
active proctype p2()
{do
  :: atomic { y++ ; y++ }
od
}
never {
do
  :: skip
  :: (x==3) -> break
od
}
```


ex6

+ Partial Order Reduction

Full statespace search for:

never claim + (never_0)
assertion violations + (if within scope of claim)
acceptance cycles - (not selected)
invalid end states - (disabled by never claim)

State-vector 36 byte, depth reached 255, ●●● errors: 0 ●●●

128 states, stored

129 states, matched

257 transitions (= stored+matched)

0 atomic steps

hash conflicts: 0 (resolved)

Stats on memory usage (in Megabytes):

0.008 equivalent memory usage for states (stored*(State-vector + overhead))

0.290 actual memory usage for states

128.000 memory used for hash table (-w24)

1.068 memory used for DFS stack (-m20000)

129.264 total actual memory usage

unreached in proctype p1

ex6.pml:6, state 7, "-end-"

(1 of 7 states)

unreached in proctype p2

ex6.pml:10, state 7, "-end-"

(1 of 7 states)

unreached in claim never_0

ex6.pml:16, state 7, "-end-"

(1 of 7 states)

pan: elapsed time 0 seconds

```
/* Smart reductions saved states */
```

```
byte x,y;
```

```
active proctype p1()
```

```
{do
```

```
  :: atomic { x++ ; x++ }
```

```
od
```

```
}
```

```
active proctype p2()
```

```
{do
```

```
  :: atomic { y++ ; y++ }
```

```
od
```

```
}
```

```
never {
```

```
do
```

```
  :: skip
```

```
  :: (x==3) -> break
```

```
od
```

```
}
```

ex7

```
/* Reductions don't work */
byte x,y;
active proctype p1()
{do
  :: atomic { x++ ; x++ }
od }
active proctype p2()
{do
  :: atomic { y++ ; y++ }
od }
never {
do
  :: skip
  :: (x==3)&&(y==2) -> break
od
}
```

ex7
do the
right depth
setting!

```
/* Reductions don't work */
byte x,y;
active proctype p1()
{do
  :: atomic { x++ ; x++ }
od }
active proctype p2()
{do
  :: atomic { y++ ; y++ }
od }
never {
do
  :: skip
  :: (x==3)&&(y==2) -> break
od
}
```

Full statespace search for:

never claim	+ (never_0)
assertion violations	+ (if within scope of claim)
acceptance cycles	- (not selected)
invalid end states	- (disabled by never claim)

State-vector 36 byte, depth reached 32767, ●●● errors: 0 ●●●

16384 states, stored

16385 states, matched

32769 transitions (= stored+matched)

0 atomic steps

hash conflicts: 3 (resolved)

Stats on memory usage (in Megabytes):

1.000	equivalent memory usage for states (stored*(State-vector + overhead))
-------	---

1.072	actual memory usage for states
-------	--------------------------------

128.000	memory used for hash table (-w24)
---------	-----------------------------------

10.681	memory used for DFS stack (-m200000)
--------	--------------------------------------

139.658	total actual memory usage
---------	---------------------------

unreached in proctype p1

ex7.pml:6, state 7, "-end-"

(1 of 7 states)

unreached in proctype p2

ex7.pml:10, state 7, "-end-"

(1 of 7 states)

unreached in claim never_0

ex7.pml:16, state 7, "-end-"

(1 of 7 states)

pan: elapsed time 0.03 seconds

pan: rate 546133.33 states/second

ex8
do the
right depth
setting!

```
/* Smart reductions saved states */
byte x,y;
active proctype p1()
{do
  :: atomic { x++ ; x++ }
od
}
active proctype p2()
{do
  :: atomic { y++ ; y++ }
od
}
never {
do
  :: skip
  :: (x==232)&&(y==2) -> break /* observe both vars to introduce
                                * state explosion
                                */
od
}
```

ex8
do the
right depth
setting!

```
i 1 p2:1 1) y = (y+1)      253
0 p1:1 1) x = (x+1)      254
Process Statement        x      y
0 p1:1 1) x = (x+1)      1      254
1 p2:1 1) y = (y+1)      2      254
spin: ex8.pml:9, Error: value (256->0 (8)) truncated in assignment
1 p2:1 1) y = (y+1)      2      255
1 p2:1 1) y = (y+1)      2      0
1 p2:1 1) y = (y+1)      2      1
1 p2:1 1) y = (y+1)      2      2
1 p2:1 1) y = (y+1)      2      3
1 p2:1 1) y = (y+1)      2      4
1 p2:1 1) v = (v+1)      2      5
```

ex9

```
byte pid1, pid2;

proctype p1()
{byte x; /* Also init to 0 */
  do
    :: x++ ; x++
  od
}
proctype p2()
{byte y; /* Also init to 0 */
  do
    :: y++ ; y++
  od
}
init {
  atomic {
    pid1 = run p1();
    pid2 = run p2();
  }
}

never {
  do
    :: skip
    :: (p1:x==2)&&(p2:y==232) -> break
  od
}
```

ex10

```
/* Channels */

chan ch = [0] of { byte }; /* rendezvous channel */

active proctype p1()
{byte x; /* local var x init to 0 */
  do
    :: x++ -> ch!x
  od
}
active proctype p2()
{byte y,z;
  do
    :: ch?y -> z++ /* z tries to keep track of the value of x */
  od
}
never {
  do
    :: skip
    :: (p2:y - p2:z) > 1 -> break
  od
}
```

ex11

```
/* Channels */

chan ch = [0] of { byte }; /* rendezvous channel */

active proctype p1()
{byte x; /* local var x init to 0 */
  do
    :: x++ -> ch!x /* ; and -> are the same */
  od
}

active proctype p2()
{byte y,z; /* can be named x, but keeping distinct names */
  do
    :: ch?y -> z++ /* z tracks value of x */
  od
}

never {
  do
    :: skip
    :: (p1:x - p2:z) > 1 -> break
  od
}
```


ex12

```
/* set depth-bound of DFS to around 20 */

chan ch = [1] of { byte }; /* buffering (non-rendezvous) channel */

active proctype p1()
{byte x; /* local var x init to 0 */
  do
    :: x++ -> ch!x /* ; and -> are the same */
  od
}
active proctype p2()
{byte y,z; /* can be named x, but keeping distinct names */
  do
    :: ch?y -> z++ /* z tracks value of x */
  od
}
never {
  do
    :: skip
    :: (p1:x - p2:z) > 2 -> break
  od
}
```

ex12b

- ex12b.pml / ex12.pml

```
/* set depth-bound of DFS to around 20 */
chan ch = [1] of { byte }; /* buffering (non-rendezvous) channel */
byte x, z;
active proctype p1()
{
  do
    :: x++ -> ch!x /* ; and -> are the same */
  od
}
active proctype p2()
{byte y; /* can be named x, but keeping distinct names */
  do
    :: ch?y -> z++ /* z tracks value of x */
  od
}
never {
  do
    :: skip
    :: (x - z) > 2 -> break
  od;
  accept: goto accept
}
```

```
warning: for p.o. reduction to be valid the never claim must be
stutter-invariant
(never claims generated from LTL formulae are stutter-invariant)
pan:1: acceptance cycle (at depth 4082)
pan: wrote ex12b.pml.trail
(Spin Version 6.4.5 -- 1 January 2016)
Warning: Search not completed
      + Partial Order Reduction
Full statespace search for:
      never claim                + (never_0)
      assertion violations       + (if within scope of claim)
      acceptance cycles         + (fairness disabled)
      invalid end states        - (disabled by never claim)
State-vector 44 byte, depth reached 6129, ●●● errors: 1 ●●●
3074 states, stored
514 states, matched
3588 transitions (= stored+matched)
0 atomic steps
hash conflicts: 0 (resolved)
Stats on memory usage (in Megabytes):
0.211      equivalent memory usage for states
(stored*(State-vector + overhead))
0.388      actual memory usage for states
128.000    memory used for hash table (-w24)
1068.115   memory used for DFS stack (-m20000000)
1196.408   total actual memory usage
pan: elapsed time 0.05 seconds
pan: rate   61480 states/second
```

ex13

```
active proctype p1()
{byte x;
  do
    :: x = x + 3 /* USER BEWARE: this statement is atomic, unlike in C !! */
  od
}
active proctype p2()
{byte y;
  do
    :: y = y + 5
  od
}
never {
  do
    :: skip
    :: (p1:x == p2:y) -> break
  od;
  accept: goto accept; /* not needed but looks Buchi */
}
```

ex13

```
active proctype p1()
{byte x;
  do
    :: x = x + 3 /* USER BEWARE: this statement is atomic, unlike in C !! */
  od
}
active proctype p2()
{byte y;
  do
    :: y = y + 5
  od
}
never {
  do
    :: skip
    :: (p1:x == p2:y) -> break
  od;
  accept: goto accept; /* not needed but looks Buchi */
}
```

<http://spinroot.com/spin/Man/Pan.html> suggests using -DNOREDUCE
PO reductions not safe with remote references (must be stutter invariant)

ex13

```
active proctype p1()
{byte x;
  do
    :: x = x + 3 /* USER BEWARE: this statement is atomic, unlike in C !! */
  od
}
active proctype p2()
{byte y;
  do
    :: y = y + 5
  od
}
never {
  do
    :: skip
    :: (p1:x == p2:y) -> break
  od;
  accept: goto accept; /* not needed but looks Buchi */
}
```

ex14

```
mtype = {are_you_free, yes, no, release}
byte progress; /* SPIN initializes all variables to 0 */
proctype phil(chan lf, rf; int philno)
{
  do
  :: do
    :: lf!are_you_free ->
      if
      :: lf?yes -> break
      :: lf?no
      fi
    od;
  do
    :: rf!are_you_free ->
      if
      :: rf?yes -> progress = 1 -> progress = 0
        -> lf!release -> rf!release -> break
      :: rf?no -> lf!release -> break
      fi
    od
  od
}
proctype fork(chan lp, rp)
{
  do
  :: rp?are_you_free -> rp!yes ->
    do
    :: lp?are_you_free -> lp!no
    :: rp?release -> break
    od
  :: lp?are_you_free -> lp!yes ->
    do
    :: rp?are_you_free -> rp!no
    :: lp?release -> break
    od
  od
}
init {
  chan c0 = [0] of { mtype }; chan c1 = [0] of { mtype };
  chan c2 = [0] of { mtype }; chan c3 = [0] of { mtype };
  chan c4 = [0] of { mtype }; chan c5 = [0] of { mtype };
  atomic {
    run phil(c5, c0, 0); run fork(c0, c1);
    run phil(c1, c2, 1); run fork(c2, c3);
    run phil(c3, c4, 2); run fork(c4, c5); }
}
never { /* Negation of []<> progress */
  do
  :: skip
  :: (!progress) -> goto accept;
  od;
  accept: (!progress) -> goto accept;
}
```

Lecture 4 : LTL model checking (CEAAT)

Question: What is the theory behind this checking?

Answer: On-the-fly LTL Model Checking using Buchi Automata!

```
- ex12b.pml / ex12.pml -
```

```
/* set depth-bound of DFS to around 20 */
chan ch = [1] of { byte }; /* buffering (non-rendezvous) channel */
byte x, z;
active proctype p1()
{
  do
    :: x++ -> ch!x /* ; and -> are the same */
  od
}
active proctype p2()
{byte y; /* can be named x, but keeping distinct names */
  do
    :: ch?y -> z++ /* z tracks value of x */
  od
}
never {
  do
    :: skip
    :: (x - z) > 2 -> break
  od;
  accept: goto accept
}
```

```
warning: for p.o. reduction to be valid the never claim must be
stutter-invariant
(never claims generated from LTL formulae are stutter-invariant)
pan:1: acceptance cycle (at depth 4082)
pan: wrote ex12b.pml.trail
(Spin Version 6.4.5 -- 1 January 2016)
Warning: Search not completed
        + Partial Order Reduction
Full statespace search for:
        never claim                + (never_0)
        assertion violations        + (if within scope of claim)
        acceptance cycles          + (fairness disabled)
        invalid end states         - (disabled by never claim)
State-vector 44 byte, depth reached 6129, ●●● errors: 1 ●●●
        3074 states, stored
        514 states, matched
        3588 transitions (= stored+matched)
        0 atomic steps
hash conflicts:          0 (resolved)
Stats on memory usage (in Megabytes):
        0.211      equivalent memory usage for states
(stored*(State-vector + overhead))
        0.388      actual memory usage for states
        128.000    memory used for hash table (-w24)
        1068.115   memory used for DFS stack (-m20000000)
        1196.408   total actual memory usage
pan: elapsed time 0.05 seconds
pan: rate      61480 states/second
```


Lecture 4 : LTL model checking (CEAAT)

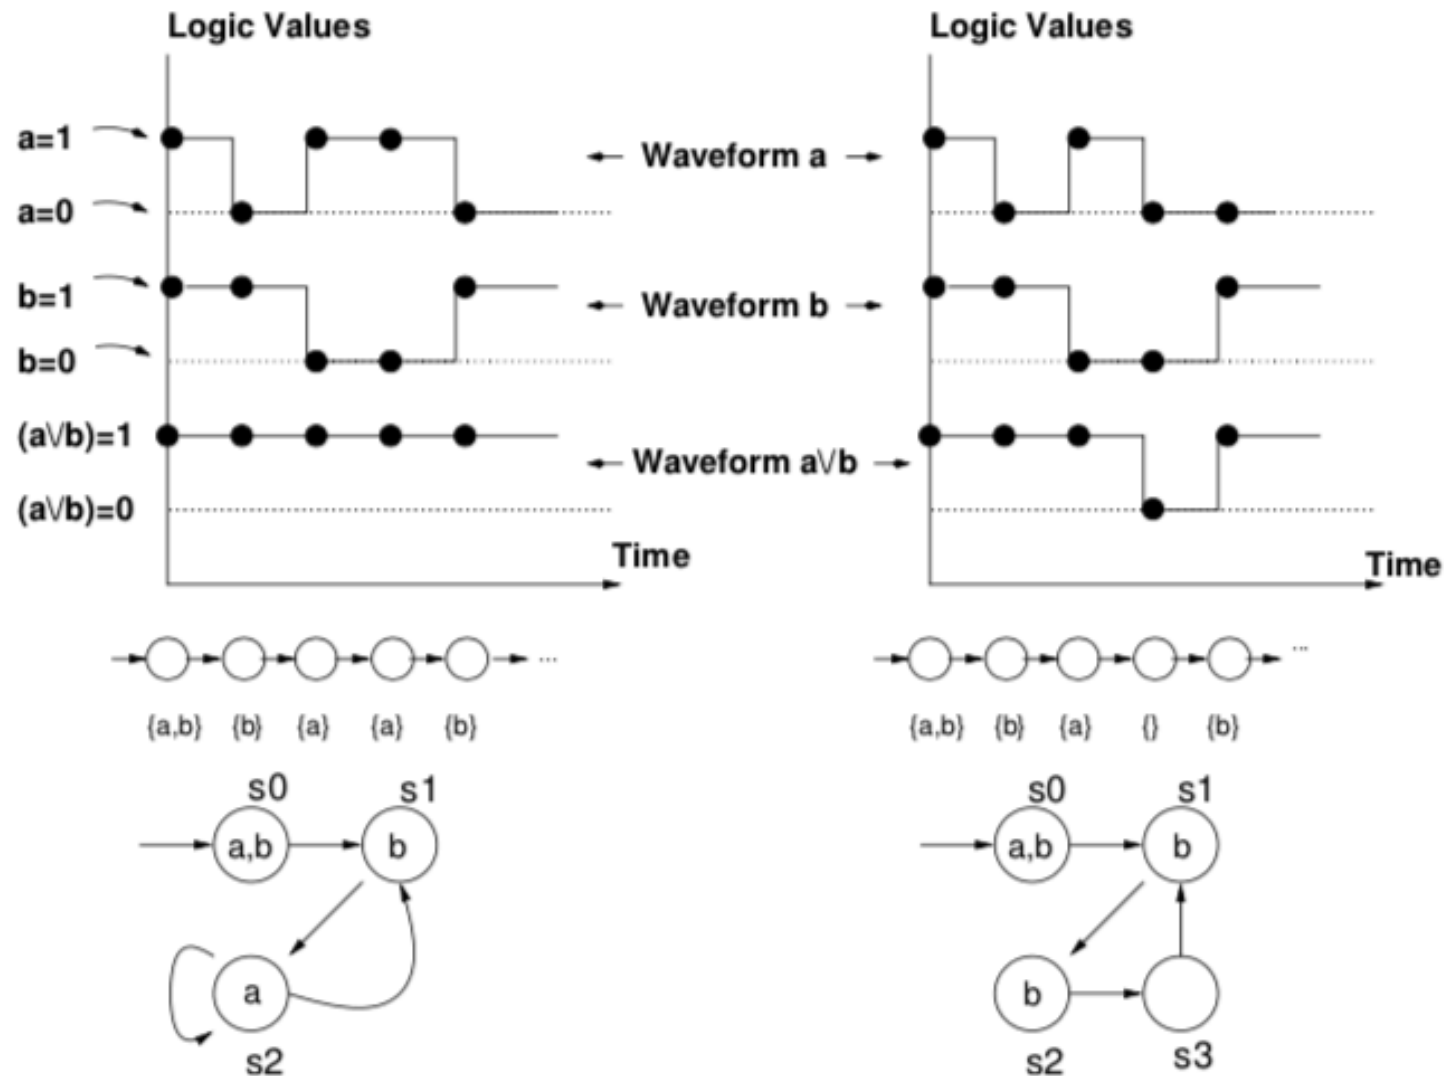


Fig. 22.1. Two Kripke structures and some of their computations. In the Kripke structure on the left, the assertion 'Henceforth ($a \vee b$)' is true.

Lecture 4 : LTL model checking (CEAAT)

22.1.5 LTL syntax

LTL formulas φ are inductively defined as follows, through a context-free grammar:

$\varphi \rightarrow x,$	a propositional variable
$\neg\varphi$	negation of an LTL formula
(φ)	parenthesization
$\varphi_1 \vee \varphi_2$	disjunction
$G\varphi$	henceforth φ
$F\varphi$	eventually φ (“future”)
$X\varphi$	next φ
$(\varphi_1 U \varphi_2)$	φ_1 until φ_2
$(\varphi_1 W \varphi_2)$	φ_1 weak-until φ_2

Lecture 4 : LTL model checking (CEAAT)

Here is the inductive definition for the semantics of LTL:

$\sigma \models x$	iff x is true at s_0 (written $s_0(x)$)
$\sigma \models \neg\varphi$	iff $\sigma \not\models \varphi$
$\sigma \models (\varphi)$	iff $\sigma \models \varphi$
$\sigma \models \varphi_1 \vee \varphi_2$	iff $\sigma \models \varphi_1 \vee \sigma \models \varphi_2$
$\sigma \models G\varphi$	iff $\sigma^i \models \varphi$ for every $i \geq 0$
$\sigma \models F\varphi$	iff $\sigma^i \models \varphi$ for some $i \geq 0$
$\sigma \models X\varphi$	iff $\sigma^1 \models \varphi$
$\sigma \models (\varphi_1 U \varphi_2)$	iff $\sigma^k \models \varphi_2$ for some $k \geq 0$ and $\sigma^j \models \varphi_1$ for all $j < k$
$\sigma \models (\varphi_1 W \varphi_2)$	iff $\sigma \models G\varphi_1 \vee \sigma \models (\varphi_1 U \varphi_2)$

Lecture 4 : LTL model checking (CEAAT)

$$\begin{array}{l} re \rightarrow \emptyset \mid \varepsilon \mid a \in \Sigma \mid (re) \mid re_1 + re_2 \mid re_1 re_2 \mid re^* \\ ore \rightarrow re^\omega \mid re \ ore \mid ore_1 + ore_2 \end{array}$$

Lecture 4 : LTL model checking (CEAAT)




Fig. 23.10. A Kripke structure (left), its corresponding Büchi automata (middle), and a property automaton expressing GFx (right)

Lecture 4 : LTL model checking (CEAAT)

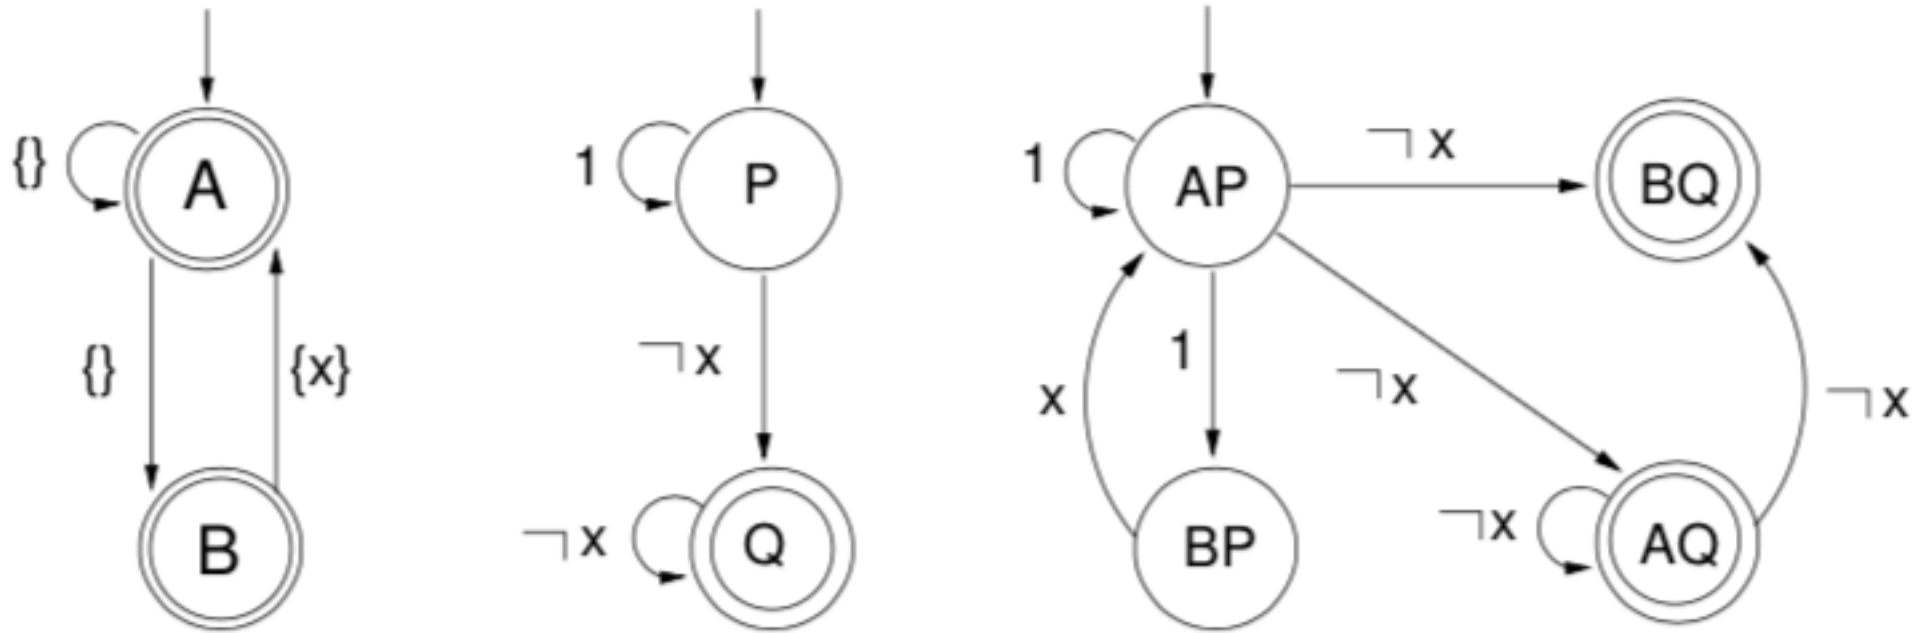


Fig. 23.11. System automaton (left), complemented property automaton (middle), and product automaton (right)

Lecture 4 : LTL model checking (CEAAT)

Notice that the “property automaton” (P) on the right-hand side of Figure 23.10 includes *all* runs that satisfy GFx . Therefore, to determine whether a given Kripke structure (“system” S) satisfies a property P , we can check whether

$$L(S) \subseteq L(P).$$


This check is equivalent to

$$L(S) \cap \overline{L(P)} = \emptyset.$$

Lecture 4 : LTL model checking (CEAAT)

$$L(S) \cap \overline{L(P)} \neq \emptyset.$$

Lecture 4 : LTL model checking (CEAAT)



S_0
S_1
S_2
S_1

S_0
S_1
S_3
S_4
S_2

S_0
S_1
S_3
S_4

S_4
S_2
S_1
S_2

S_4
S_2
S_1
S_3
S_4